

[Why
CoreWeave](#)[Pricing](#)[Documentation](#)[Company](#) ▼[Contact
Us](#)[Login](#)

CoreWeave Cloud Pricing

CoreWeave's pricing is designed for flexibility. Instances are highly configurable, giving you the freedom to customize GPU, CPU, RAM, and storage requests when scheduling your workloads.

Our entire infrastructure is purpose-built for compute-intensive workloads, and everything from our servers to our storage and networking solutions are designed to deliver best-in-class performance that are up to **35× faster and 80% less expensive** than generalized public clouds.

CoreWeave GPU Cloud Pricing

CoreWeave Cloud GPU instance pricing is highly flexible, and meant to provide you with ultimate control over configuration and cost. Pricing below is a la carte, where the total instance cost is a combination of a GPU component, the number of vCPU, and the amount of RAM allocated. To keep things simple, CPU and RAM cost are the same per base unit, and the only variable is the GPU chosen for your workload or Virtual Server.

A valid GPU instance configuration must include at least 1 GPU, at least 1 vCPU and at least 2GB of RAM. When deploying a Virtual Server, the GPU instance configuration must also include at least 40GB of root disk NVMe tier storage.

GPU Model	VRAM (GB)	Max vCPUs per GPU (\$0.01/hr)	Max RAM (GB) per GPU (\$0.005/hr)	GPU Component Cost Per Hour
-----------	-----------	-------------------------------------	---	-----------------------------------


[Why
CoreWeave](#)
[Pricing](#)
[Documentation](#)
[Company ▾](#)
[Contact
Us](#)
[Login](#)

A100 80GB NVLINK	80	48	256	\$2.21
A100 80GB PCIe	80	48	256	\$2.21
A100 40GB NVLINK	40	48	256	\$2.06
A100 40GB PCIe	40	48	256	\$2.06
A40	48	48	256	\$1.28
RTX A6000	48	48	256	\$1.28
RTX A5000	24	36	128	\$0.77
RTX A4000	16	36	128	\$0.61
Quadro RTX 5000	16	36	128	\$0.57
Quadro RTX 4000	8	36	128	\$0.24
Tesla V100 NVLINK	16	36	128	\$0.80

*Compute instances on CoreWeave Cloud are configurable. For more info, please refer to our [Resource Based Pricing Documentation](#).



Why
CoreWeave

Pricing

Documentation

Company ▾

Contact
Us

Login

CPU Model	RAM per vCPU	Cost Per vCPU
AMD EPYC Milan	4	\$0.035
AMD EPYC Rome	4	\$0.03
Intel Xeon Ice Lake	4	\$0.035
Intel Xeon Scalable	4	\$0.03
Intel Xeon v4	4	\$0.02
Intel Xeon v3	4	\$0.0125

*Compute instances on CoreWeave Cloud are configurable. For more info, please refer to our [Resource Based Pricing Documentation](#).

Storage

CoreWeave Cloud’s high-performance network attached storage is priced to be simple to understand and budget, with no limitations on scale, throughput or IOPS. Without extraneous charges for transferring

Networking

Networking on CoreWeave is built to empower scale computational workloads. We don’t charge for ingress and egress, and if your workload requires significant throughput we will work with our partners to



Why
CoreWeave

Pricing

Documentation

Company ▾

Contact
Us

Login

Block Storage	HDD	\$0.04 per GB / month	Public IP Address	\$4.00 per IP / month
	NVMe	\$0.07 per GB / month		
File System Storage	Standard HDD	\$0.04 per GB / month	Burstable IP Address	\$10.00 per IP / month
	Standard NVMe	\$0.07 per GB / month		
	Premium NVMe	\$0.10 per GB / month	VPC	\$20.00 per VPC / month
Object Storage	HDD and NVMe	\$0.03 per GB / month		

GPU Cloud Pricing FAQ

How is CoreWeave up to 80% less expensive?

[Why
CoreWeave](#)[Pricing](#)[Documentation](#)[Company ▾](#)[Contact
Us](#)[Login](#)

Do you offer discounts?

The vast majority of our clients enjoy the benefits of CoreWeave's on-demand pricing with no contract or commitment. For workloads with relatively predictable usage patterns, we offer two types of volumetric discounts: Reserved Instances (up to 60% for 24/7 committed usage) and Bulk Credits (up to 25% for large-scale, burst compute use cases consumed on-demand). Reach out to our Sales team to learn more about how these discounts can help your business.

Do you charge for bandwidth/data transfer?

No! Gross. Other cloud providers "tax" you for each GB you transfer, making the cost to run at scale in the cloud prohibitive. At CoreWeave, we don't charge for things like region-to-region transfers, workstation data, or egress in the vast majority of use cases. If your bandwidth needs are extraordinary, we'll work to understand your needs before negotiating flat bandwidth rates with our suppliers on your behalf.


Do you offer "spot" or interruptible instances?

No, we do not. You can think of our on-demand pricing as "interruptible pricing for uninterruptible instances". Our clients benefit from economics that allow them to scale, with peace of mind knowing that when they spin-up instances on CoreWeave Cloud, it's theirs until they choose to spin it down once their workloads are complete.

[Why
CoreWeave](#)[Pricing](#)[Documentation](#)[Company](#) ▾[Contact
Us](#)[Login](#)

deliver economics that scale for years to come.

RECENT BLOG POSTS



CoreWeave Named to USA Today's '2024 Top Workplaces'

[Read more →](#)

The Redesign of the Data Center Has Already Started. Here's What It Looks Like

[Read more →](#)

Mistral AI and CoreWeave Demonstrate Partnership at NVIDIA GTC, Mistral AI Hackathon

[Read more →](#)

[Why
CoreWeave](#)[Pricing](#)[Documentation](#)[Company](#) ▼[Contact
Us](#)[Login](#)

CoreWeave is a specialized cloud provider, delivering a massive scale of GPUs on top of the industry's fastest and most flexible infrastructure.

101 Eisenhower Pkwy Ste 106

Roseland NJ 07068-1050

Why CoreWeave **Careers**

Documentation **Blog**

Pricing **Events & Webinars**

Contact Sales **Login**

About Us **Status**

Support **Terms of Service**



Why
CoreWeave

Pricing

Documentation

Company ▾

**Contact
Us**

Login