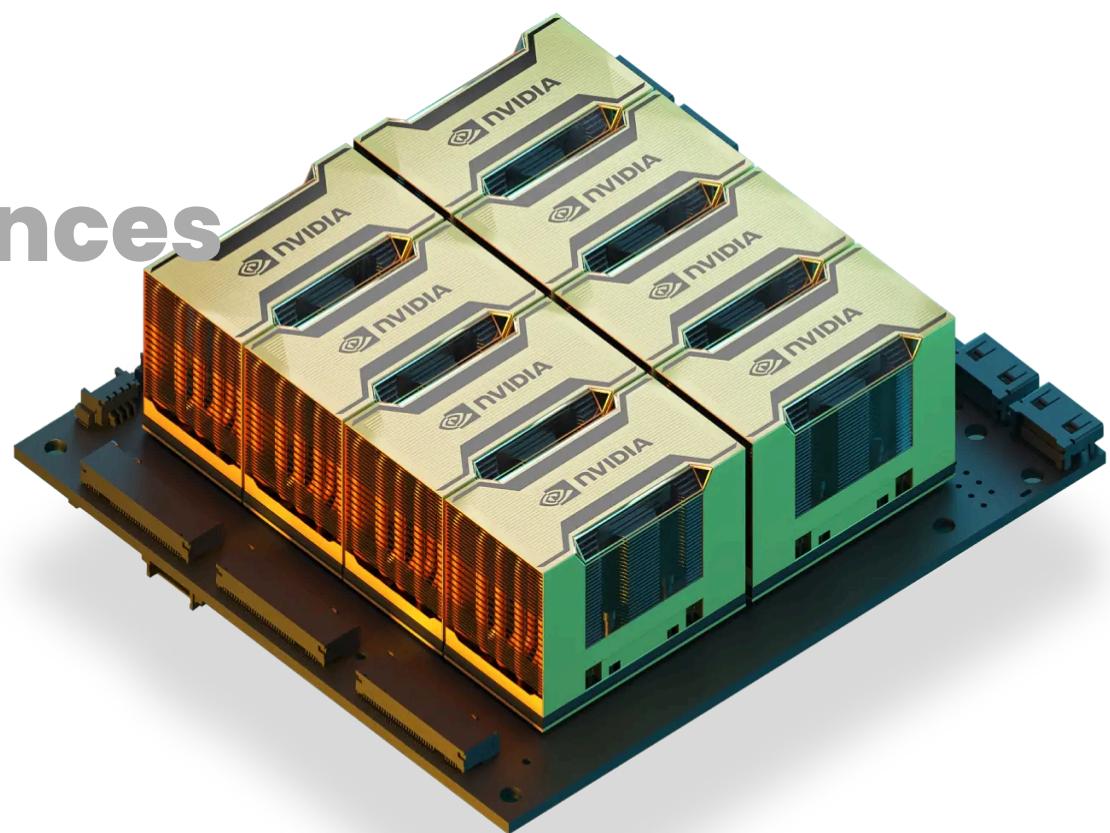


NVidia® Instances

Simple & clear. Easy to setup.

[Get Started](#)

*All prices are in US dollars.



The H100 virtual dedicated servers are powered by:

Up to 8 NVidia® H100 80GB GPUs, each containing 16896 CUDA cores and 528 Tensor Cores.

This is the current flagship silicon from NVidia®, unbeaten in raw performance for AI operations.

We deploy the SXM5 NVLINK module, which offers a memory bandwidth of 2.6 Gbps and up to 900GB/s P2P bandwidth.

Fourth generation AMD Genoa, up to 384 threads with a boost clock of 3.7GHz

The name 8H100.80S.360V is composed as follows: 8x H100 SMX5, 360 CPU core threads & virtualized.

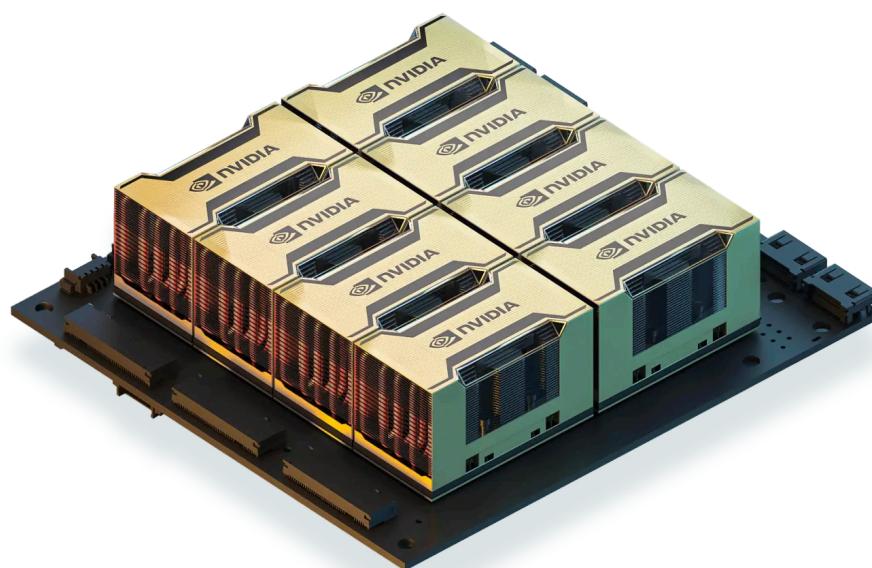


INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8H100.80S.176V	H100 SXM5 80GB	8	176	1480	640	/	\$20.72/h	\$19.89/h	\$15.54/h
4H100.80S.120V	H100 SXM5 80GB	4	120	480	320	/	\$10.36/h	\$9.95/h	\$7.77/h
4H100.80S.88V	H100 SXM5 80GB	4	88	740	320	/	\$10.36/h	\$9.95/h	\$7.77/h
2H100.80S.60V	H100 SXM5 80GB	2	60	240	160	/	\$5.18/h	\$4.97/h	\$3.88/h
1H100.80S.30V	H100 SXM5 80GB	1	30	120	80	/	\$2.59/h	\$2.49/h	\$1.94/h

Name	Model							Demand	Price	Price	
Instance Name	GPU Model	GPU	CPU	RAM	VRAM	P2P	Price On Demand	6-Month Price	2-Year Price	4H100.	
Instance Name	GPU Model	GPU	CPU	RAM	VRAM	P2P	Price On Demand	6-Month Price	2-Year Price	2H100.	
Instance Name	GPU Model	GPU	CPU	RAM	VRAM	P2P	Price On Demand	6-Month Price	2-Year Price	1H100.	
◀ ▶											

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

The A100 virtual dedicated servers are powered by:



Up to 8 NVidia® A100 80GB GPUs, each containing 6912 CUDA cores and 432 Tensor Cores.

We only use the SXM4 'for NVLINK' module, which offers a memory bandwidth of over 2TB/s and Up to 600GB/s P2P bandwidth.

Second generation AMD EPYC Rome, up to 192 threads with a boost clock of 3.3GHz.

The name 8A100.176V is composed as follows: 8x RTX A100, 176 CPU core threads & virtualized.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8A100.40S.176V	A100 SXM4 40GB	8	176	960	320	/	\$14.00/h	\$13.44/h	\$10.50/h
4A100.40S.88V	A100 SXM4 40GB	4	88	480	160	/	\$7.00/h	\$6.72/h	\$5.25/h
2A100.40S.44V	A100 SXM4 40GB	2	44	240	80	/	\$3.50/h	\$3.36/h	\$2.63/h
1A100.40S.22V	A100 SXM4 40GB	1	22	120	40	/	\$1.75/h	\$1.68/h	\$1.31/h
8A100.176V	A100 SXM4 80GB	8	176	960	640	600 GB/s	\$15.61/h	\$14.99/h	\$11.71/h
4A100.88V	A100 SXM4 80GB	4	88	480	320	300 GB/s	\$7.80/h	\$7.49/h	\$5.85/h
2A100.44V	A100 SXM4 80GB	2	44	240	160	100 GB/s	\$3.90/h	\$3.74/h	\$2.92/h
1A100.22V	A100 SXM4 80GB	1	22	120	80	/	\$1.95/h	\$1.87/h	\$1.46/h

NAME	MODEL							DEMAND	PRICE	PRICE
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2A100.4
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1A100.4
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	8A100.4
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	4A10
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2A10
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1A10

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

The L40S virtual dedicated servers are powered by:

Up to 8 Nvidia® L40S 48GB GPUs, each containing 18176 CUDA cores and 568 Tensor Cores.

Features 864GB/s memory bandwidth.

Tensor Performance: 1466 Tflops

RT Core Performance: 212 Tflops

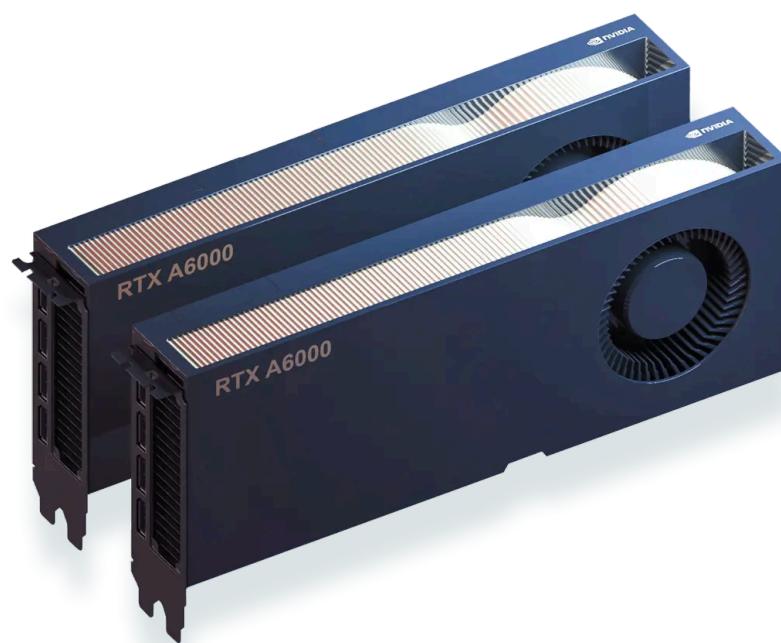
Single-precision performance: 91.6 Tflops

The name 8L40S.160V is composed as follows: 8x L40S, 160 CPU core threads & virtualized.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8L40S.160V	NVidia L40S	8	160	480	384	/	\$10.86/h	\$10.43/h	\$8.14/h
4L40S.80V	NVidia L40S	4	80	240	192	/	\$5.43/h	\$5.21/h	\$4.07/h
2L40S.40V	NVidia L40S	2	40	120	96	/	\$2.72/h	\$2.61/h	\$2.04/h
1L40S.20V	NVidia L40S	1	20	60	48	/	\$1.36/h	\$1.31/h	\$1.02/h

NAME	MODEL							DEMAND	PRICE	PRICE
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2L40S.
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1L40S.

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.



The RTX A6000 virtual dedicated servers are powered by:

Up to 8 NVidia® RTX A6000 [2021] GPUs, each containing 10752 CUDA cores, 336 Tensor Cores and 84RT cores.

Despite having less tensor cores than the V100, it is able to process tensor operations faster due to a different architecture.

Second generation AMD EPYC Rome, up to 96 threads with a boost clock of 3.35GHz.

PCIe Gen4 for faster interactions between GPUs.

The name 8A6000.80V is composed as follows: 8x RTX A6000, 80 CPU core threads & virtualized.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8A6000.80V	NVidia RTX A6000 48GB	8	80	480	384	/	\$8.06/h	\$7.74/h	\$6.04/h
4A6000.40V	NVidia RTX A6000 48GB	4	40	240	192	/	\$4.03/h	\$3.87/h	\$3.02/h
2A6000.20V	NVidia RTX A6000 48GB	2	20	120	96	/	\$2.01/h	\$1.93/h	\$1.51/h
1A6000.10V	NVidia RTX A6000 48GB	1	10	60	48	/	\$1.01/h	\$0.97/h	\$0.76/h

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	8A6000
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	4A6000
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2A6000
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1A6000

The RTX6000 ADA virtual dedicated servers are powered by:

Up to 8 NVidia® RTX6000 ADA [2022] GPUs, each containing 18176 CUDA cores, 568 fourth-generation Tensor Cores and 142RT cores.

Features 48GB of GDDR6 memory for working with the largest 3D models, render images, simulation and AI datasets.

The name 8RTX6000ADA.80V is composed as follows: 8x RTX6000ADA, 80 CPU core threads & virtualized.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8RTX6000ADA.80V	NVidia RTX6000 Ada 48GB	8	80	480	384	/	\$10.18/h	\$9.77/h	\$7.63/h
4RTX6000ADA.40V	NVidia RTX6000 Ada 48GB	4	40	240	192	/	\$5.09/h	\$4.89/h	\$3.82/h
2RTX6000ADA.20V	NVidia RTX6000 Ada 48GB	2	20	120	96	/	\$2.55/h	\$2.45/h	\$1.91/h
1RTX6000ADA.10V	NVidia RTX6000 Ada 48GB	1	10	60	48	/	\$1.27/h	\$1.22/h	\$0.95/h

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	8RTX60
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	4RTX60
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2RTX60
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1RTX60

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

The V100 virtual dedicated servers are



Second generation Xeon Scalable 4214R CPUs [2020], up to 48 threads with a boost clock of 3.5GHz.

NVLink for high bandwidth P2P communication.

The name 4V100.20V is composed as follows: 4x V100, 20 CPU core threads & virtualized.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
8V100.48V	NVidia Tesla V100 16GB	8	48	180	128	NVLink up to 50GB/s	\$5.30/h	\$5.09/h	\$3.97/h
4V100.20V	NVidia Tesla V100 16GB	4	20	90	64	NVLink up to 50GB/s	\$2.65/h	\$2.54/h	\$1.99/h
2V100.10V	NVidia Tesla V100 16GB	2	10	45	32	NVLink up to 50GB/s	\$1.32/h	\$1.27/h	\$0.99/h
1V100.6V	NVidia Tesla V100 16GB	1	6	23	16	/	\$0.66/h	\$0.63/h	\$0.49/h

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	8V100.4
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	4V100.2
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2V100.1
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1V100.0



*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

The AMD 7900XTX virtual dedicated servers are powered by:

AMD's flagship GPU, offering an economical path for ML and AI workloads.

With 24 GB of GDDR6 memory and advanced RDNA 3 architecture, it's tailor-made for demanding ML workloads and experimenting with new AI models.

Ideal for AI enthusiasts and professionals, we prioritize users engaged in AI-related projects.

4RX7900XTX.40V	AMD 7900XTX	4	40	160	96	/	\$1.56/h	\$1.50/h	\$1.17/h
2RX7900XTX.20V	AMD 7900XTX	2	20	80	48	/	\$0.78/h	\$0.75/h	\$0.58/h
1RX7900XTX.10V	AMD 7900XTX	1	10	40	24	/	\$0.39/h	\$0.37/h	\$0.29/h

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	12RX79
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	8RX79
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	4RX79
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	2RX79
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	1RX79



*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.



The CPU virtual dedicated servers are powered by:

Second or Third generation AMD EPYC Rome or Milan.

All hardware is dedicated to your server for the best performance.

The name CPU.32V indicates the server runs on 32 virtualized core threads.

INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE
CPU.120V.480G	AMD EPYC	0	120	480	0	/	\$4.44/h	\$4.26/h	\$3.33/h
CPU.96V.384G	AMD EPYC	0	96	384	0	/	\$3.55/h	\$3.41/h	\$2.66/h
CPU.64V.256G	AMD EPYC	0	64	256	0	/	\$2.37/h	\$2.28/h	\$1.78/h
CPU.32V.128G	AMD EPYC	0	32	128	0	/	\$1.18/h	\$1.13/h	\$0.89/h
CPU.16V.64G	AMD EPYC	0	16	64	0	/	\$0.59/h	\$0.57/h	\$0.44/h
CPU.8V.32G	AMD EPYC	0	8	32	0	/	\$0.30/h	\$0.29/h	\$0.22/h
CPU.4V.16G	AMD EPYC	0	4	16	0	/	\$0.15/h	\$0.14/h	\$0.11/h

NAME	MODEL							DEMAND	PRICE	PRICE	
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	CPU.64	
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	CPU.32	
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	CPU.16	
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	CPU.8	
INSTANCE NAME	GPU MODEL	GPU	CPU	RAM	VRAM	P2P	PRICE ON DEMAND	6-MONTH PRICE	2-YEAR PRICE	CPU.4	

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

Storage

Our instances run on a network storage cluster. This cluster allows us to constantly keep your data in 3 copies, to ensure redundancy in the event of hardware failure.

Our NVME cluster offers high IOPS and excellent continuous bandwidth, the HDD cluster is ideal for larger datasets. By default, the volume sizes are limited, however, the limits can be increased on demand.

Products	Company	API	Docs
GPU Instances	Contact Us		
GPU Clusters	Jobs		
Serverless Inference	Career		
ML Pipeline Builder	Terms of Service		
	Privacy Policy		

Type	Continuous Bandwidth [MB/s]	Burst Bandwidth [MB/s]	IOPS	Internal Network Speed [Gbit/s]	Price [\$/GB/month]	HDD	250	2000	300	50

*Note: Once you deploy your instance, the price and discount are fixed and not subject to future changes.

datacenter facility
and are owned and
operated solely by
DataCrunch. All
DataCrunch
operations are
ISO27001
certified.

secured by storing
it in 3 copies at all
times.

allocated to your
machine and your
machine only.